



ENGR-UH 3511

Computer Organization and Architecture

Name: Nishant Aswani

Net ID: nsa325

Assignment Title: Final Project Report

Compiling a LSTM Neural Network for Various ISAs using the Glow and LLC Compilers

Nishant Aswani, nsa325@nyu.edu

Computer Organization and Architecture (ENGR-UH 3511), Professor Mihalios Maniatakos

1 Introduction

With the advent of machine learning (ML) techniques being applied for forecasting and vision, along with increasing abilities to gather data, rose the need for specific hardware. Traditional CPUs are unable to efficiently exploit the parallelism in the operations carried out in neural networks [1]. Application-specific integration circuits (ASICs) for machine learning, on the other hand, are designed specifically to exploit the repetitive and predictable operations of a neural network. Some of the main principles of domain-specific architectures (DSAs) include "dedicated local memories, large numbers of arithmetic units, [and] simple forms of parallelism", in the hopes of improving energy efficiency [2].

As a result, there is a need for compilers capable of optimizing deep neural networks (DNN) for execution on DSAs. Traditionally, frameworks "iterate over the nodes in the graph and execute them one by one" [2], leading to immense inefficiencies. Compilers, such as Glow, take into account the entire network representation before making compilation decision, in hopes to improve efficiency regardless of hardware.

This project makes use of Glow to compile a bidirectional Long short-term memory (LSTM) model on various instruction set architectures (ISAs) in hopes to explore what typical ML operations look like as instructions as well as compare how relatively complex instructions are implemented in different architectures.

1.1 LSTM Model

An LSTM is a form of recurrent neural network (RNN) designed as a solution to the "vanishing gradient" [3] problem that RNN models face: context remembrance across long time lags "takes a prohibitive amount of time or does not work at all" [3]. LSTMs have additional gates (forget, input, and output gates) compared to traditional RNNs which allow for them to be more robust and aid the model in deciding whether certain values should be used or not.

What makes an LSTM layer interesting is that it carries out more complex operations than the typical Dense layer of neurons, where a multiply and accumulate (MAC) is the dominant operation. Furthermore, looking at figure 1, we see within the gates (yellow containers) the use of sigmoid and tanh functions as activation functions for the LSTM cell.





Figure 1: Overview of a LSTM cell

1.2 The `tanh` function

As implied above, the `tanh` function is a popular motif in machine learning models. The purpose of an activation function is to build a model to explain data which is non-linear. Without activation functions, neural networks would collapse to linear approximations.

In hardware, we must find a numerical method to compute such functions. One of the models for the `tanh` function is shown here [4]:

Figure 2: Computing the $\tanh(x)$ function

2 Methodology

The model used in this project, shown in Figure 3, was obtained in the ONNX file format from a unit test of Microsoft's NN scoring engine and visualized using Netron.

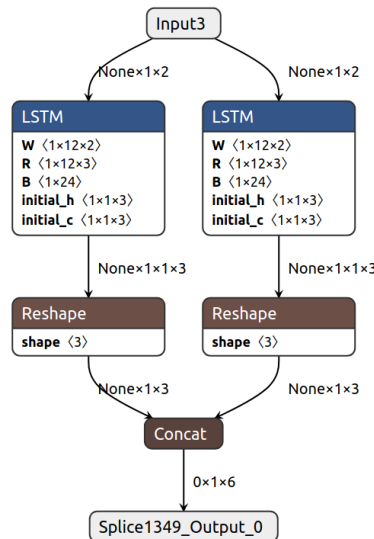


Figure 3: Visualizing the LSTM model

The Glow compiler was built in release mode, which provided a `model-compiler` binary capable of handling ONNX models as an input.

Using the following shell command, a '.s' file containing assembly code was produced.:

```
./model-compiler \
-model="../../bilstm.onnx" -emit-bundle="../../mips32_bundle" \
-backend=CPU -onnx-define-symbol=None,1 \
-llvm-compiler=/usr/bin/llc-8 -llvm-compiler-opt=-filetype=asm \
-target=mips -mcpu=mips32 \
-dump-graph-DAG=graph.dot -dump-llvm-ir > ir.txt
```

The `-model` flag takes the path of the ONNX model, `-emit-bundle` is the output directory. The `-llvm-compiler` flag accepts an external compiler (in place of the default llvm compiler), for which the filetype was set to `asm` to obtain assembly code instead of machine code. The `-target=mips -mcpu=mips32` flags tell the llc compiler which architecture to build for. Finally, the `-dump` flags output the graph representation of the network and the low-level intermediate representation (IR), respectively.

Figure 5 shows a slice of the the graph representation of the neural network. Towards the middle we see the activation functions in their own nodes. Hence, as part of the exploration, it was sought to discover how activation nodes were represented in the low-level IR and assembly language.

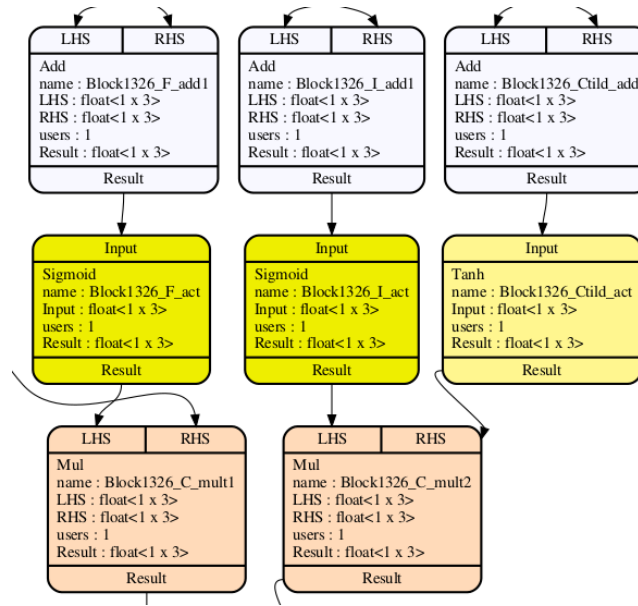


Figure 4: Graph representation using Glow

3 Results and Discussion

3.1 Intermediate Representation

Taking a look at the intermediate representation, multiple loops were found defined as **stacked kernels**. Glow "stacks operators and performs a few data-parallel operators one after the other on the same memory location,"[2] an optimization that cannot allegedly done by a traditional LLVM compiler.

The loop below represents the right hand side of Figure 5:

```

loop:                                     ; preds = %loop, %entry
    %3 = phi i64 [ 0, %entry ], [ %nextvar, %loop ]
    %4 = call float @libjit_element_add_kernel_f(i64 %3, float* %0, float* %1, float* null)
    %buffer.element.addr = getelementptr float, float* %0, i64 %3
    store float %4, float* %buffer.element.addr
    %5 = call float @libjit_tanh_kernel_f(i64 %3, float* %0, float* null, float* null)
    %buffer.element.addr1 = getelementptr float, float* %0, i64 %3
    store float %5, float* %buffer.element.addr1
    %6 = call float @libjit_element_mul_kernel_f(i64 %3, float* %2, float* %0, float* null)
    %buffer.element.addr2 = getelementptr float, float* %2, i64 %3
    store float %6, float* %buffer.element.addr2
    %nextvar = add nuw nsw i64 %3, 1
    %loopcond = icmp ult i64 %nextvar, 3
    br i1 %loopcond, label %loop, label %afterloop, !llvm.loop !6

afterloop:                               ; preds = %loop
    ret void
}

```

We see that three operators `add`, `tanh`, and `mul` have been combined into a single stack (see variables %4, %5, %6). Furthermore, to represent the `tanh` function there is a call to the `libjit tanh kernel f` function. This tag will prove useful when parsing the assembly code for the different architectures.

The source files of the Glow compiler reveal how the `tanh` function is implemented.

```

DEFINE_DATA_PARALLEL_KERNEL(libjit_tanh_kernel_f, float,
                             1 - 2 / (expf(LHS[idx] * 2) + 1))

```

Unsurprisingly, Glow uses the most accurate way to compute the `tanh` function. This function is also seen in figure 2.

3.2 Compiling for MIPS32

After compiling for the MIPS ISA, there were multiple `jal` instructions such as these:

```

jal    libjit_stacked_kernel.1_3_specialized

```

These labels were defined below and corresponded to assembly instructions of what was seen in the the intermediate representation. **For MIPS, the stacked kernel 1 and 3 were 60 lines of assembly instructions.** Since the label implies that this is a set of instructions for two stacked kernels, it was unclear which of the three `expf` operations belonged to the `tanh` function. Although, it is quite likely that the following chunk of instructions carries the `tanh` operation:

```

mul.s    $f0, $f1, $f0
swc1     $f0, 0($16)
lwc1     $f0, 4($17)
lwc1     $f1, 4($18)
add.s    $f0, $f1, $f0
jal      expf
add.s    $f12, $f0, $f0
add.s    $f0, $f0, $f20
div.s    $f0, $f21, $f0
sub.s    $f0, $f20, $f0

```

Interestingly, the MIPS assembly was littered with `4-byte Folded Spill` comments generated by the compiler. Given that these comments were beside `sw` instructions, it is likely that the 32 registers of the MIPS CPU were not enough for the computations carried out. Hence, the compiler is forced to interact with the memory, which is an expensive instruction. This provides an emphasis for why DSAs exist. For machine learning applications, models need to produce results instantaneously. Larger memories on the ASICs reduce memory interactions, in turn reducing communication overhead.

3.3 Compiling for ARM Cortex A-72

The ARM Cortex A-72 chip is used onboard the Raspberry Pi 4. **For assembly with ARM, the stacked kernel 1 and 3 were 64 lines of assembly instructions.**

Regarding stack pointer manipulation, ARM had a clear advantage over the MIPS ISA. The first 5 instructions in the MIPS file were used to manipulate stack pointers, while ARM used a single `push` instruction.

The same operations shown above in MIPS were found in the ARM assembly:

```
bl      __mulsf3
str     r0, [r4]
ldr     r0, [r5, #4]
ldr     r1, [r6, #4]
bl      __addsf3
mov     r1, r0
bl      __addsf3
bl      expf
mov     r1, #1065353216
bl      __addsf3
mov     r1, r0
mov     r0, #1073741824
bl      __divsf3
mov     r1, r0
mov     r0, #1065353216
bl      __subsf3
```

Clearly, the same operation takes more ARM instructions than MIPS instructions. It was unclear whether having to use the `mov` instruction after operations is an ARM limitation or a missed optimization.

Despite not decreasing the quantity of instructions, the ARM assembly code used far fewer load/store instructions, making less costly choices. Furthermore, there were no spills in the ARM instructions.

3.4 Compiling for x86 Skylake-AVX512 with AVX512NNI Enabled

AVX-512 was a set of instruction extensions to the SIMD architecture with several new features. Namely, the width of the SIMD registers increased from 128 to 256 bits and introduced fast multiply and accumulate (FMA) instructions. For this compilation, the AVX512NNI feature was enabled, which adds specific vector instructions for deep learning.

One of the two instructions added is the `VPDPBUSD` instruction capable of speeding up the multiply and accumulate operation. Figure 5 shows a representation of this.

However, it was extremely disappointing to find that none of the AVX512 instructions had been used in the assembly code produced. It seemed as if the model had not obeyed the architecture specification flags, despite producing x86 assembly code. Moreover, the **stacked kernel 1 3** operation was carried out as it had been in the previous runs, just with x86-specific instructions. **For this implementation, the stacked kernel 1 and 3 were 53 lines of assembly instructions.**

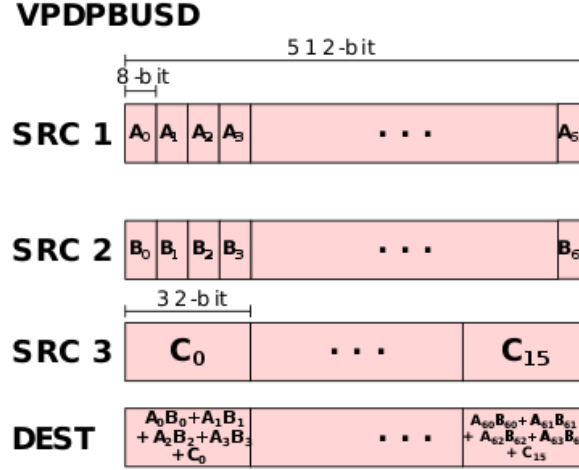


Figure 5: Visual explanation of the VPDPBUSD instruction

The supposed tanh implementation is shown below; it undeniably uses fewer instructions.

```
mulss    (%r14), %xmm0
movss    %xmm0, (%r14)
movss    4(%rbx), %xmm0          # xmm0 = mem[0],zero,zero,zero
addss    4(%r15), %xmm0
addss    %xmm0, %xmm0
callq    expf
movss    .LCPI4_0(%rip), %xmm2   # xmm2 = mem[0],zero,zero,zero
addss    %xmm2, %xmm0
movss    .LCPI4_1(%rip), %xmm1   # xmm1 = mem[0],zero,zero,zero
divss    %xmm0, %xmm1
movaps    %xmm2, %xmm0
subss    %xmm1, %xmm0
```

It was thought that perhaps the LSTM model being used had little to no room to employ the AVX instructions. So, for thoroughness, a dense network (which heavily consist of MAC operations) was compiled as well. The new model compiled to use some of the AVX512 instructions, such as `vpbroadcastb`. Despite having multiple locations for MAC instructions, the model still did not use any `vpmadd52huq` instructions. Although this attempt at compilation could not exploit DLP with the AVX512 instructions, the use of some vector instructions still reduced the total instruction count significantly. Obviously, there were no cases of spillovers, unlike MIPS.

4 Conclusion

As a result of this exploration, a basic understanding of the purpose and design of hardware accelerators was gained. Specifications such as on-board memory or high-bandwidth communication are crucial to producing adequate performance for a machine learning ASIC. Furthermore, the need for a heterogeneous machine learning compiler was also understood. Glow is able to optimize the machine learning model before passing it on for cross-compilation. Exploring the results in different architectures and CPUs was an interesting exercise, showing the limitations and benefits of various architectures. For example, the MIPS CPU's limitation of 32 registers leads to immense communication overhead with memory.

Habana, a hardware accelerator startup now acquired by Intel, integrated the Glow compiler as part of their backend. Habana built Goya, a "VLIW CPU based on its own instruction-set architecture, calling it the TPC (Tensor Processor Core)" [5] (see Figure 6). Their chip was able perform quite well in the benchmarks, surpassing NVIDIA's Tesla T4. It is hardware accelerators such as these that benefit from opensource software like Glow, producing superior products while learning and giving back to the ML community.

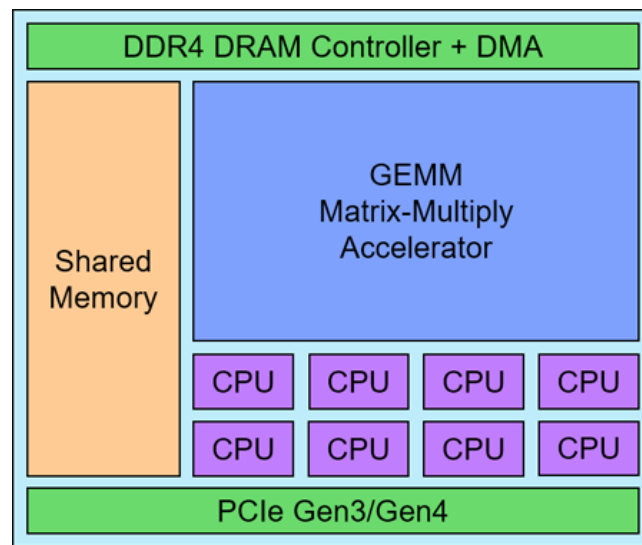


Figure 6: The Habana chip

References

- [1] Vivienne Sze et al. "Hardware for machine learning: Challenges and opportunities". In: *2017 IEEE Custom Integrated Circuits Conference (CICC)*. IEEE. 2017, pp. 1–8.
- [2] Nadav Rotem et al. "Glow: Graph lowering compiler techniques for neural networks". In: *arXiv preprint arXiv:1805.00907* (2018).
- [3] Sepp Hochreiter and Jürgen Schmidhuber. "Long Short-Term Memory". In: *Neural Computation* 9.8 (1997), pp. 1735–1780. DOI: 10.1162/neco.1997.9.8.1735. eprint: <https://doi.org/10.1162/neco.1997.9.8.1735>. URL: <https://doi.org/10.1162/neco.1997.9.8.1735>.
- [4] Nelson HF Beebe. "Accurate hyperbolic tangent computation". In: *Technical report, Center for Scientific Computing, Department of Mathematics, University of Utah* (1991).
- [5] Linley Gwennap. "HABANA WINS CIGAR FOR AIINFERENCE". In: (2019).